



ThinkStack: An Offline Small-Language-Model Research Assistant for Literature Review and Writing

by

Aditya Mehta (2440204)
Jitvan Chadha (2440222)
Ritesh K R (2440233)

Under the guidance of

Dr. New Begin

A Project report submitted in partial fulfilment of the requirements of
V Semester B.Sc. (Computer Science and Statistics)

CHRIST (Deemed to be University)

August 2026



CERTIFICATE

This is to certify that the report titled **ThinkStack: An Offline Small-Language-Model Research Assistant for Literature Review and Writing** is a bona fide record of work done by **Aditya Mehta (2440204)**, **Jitvan Chadha (2440222)** and **Ritesh K R (2440233)** of CHRIST (Deemed to be University), Bangalore, in partial fulfilment of the requirements of V Semester B.Sc. (Computer Science and Statistics) during the year 2026–2027.

Head of the Department

Project Guide

Valued-by:

1. _____

2. _____

Examination Centre : CHRIST (Deemed to be University)

Date of Exam : _____

DECLARATION OF AI TOOL USAGE

We declare that the following artificial-intelligence tools were used during the development of this project, in the roles stated.

The design of the software: its architecture, the decision to run every inference on the user's own machine, the module boundaries, the data model and the behaviour of each feature: is the developers' own. Claude was used to implement the front end and to draft documentation. Every contribution it made was reviewed before it was accepted.

Tool	Role in the project
Claude (Anthropic)	Implementation of the front end, and drafting of documentation. Software design, architecture and feature behaviour were specified by the developers; the decision log in <code>docs/ADR.md</code> records that reasoning, and every generated change was reviewed before it was committed.
Qwen2.5 0.5B / 1.5B (Alibaba, quantised GGUF)	Not a development tool but the <i>product itself</i> . These models are bundled with, or downloaded by, the application and perform summarisation, claim extraction, gap analysis and L ^A T _E X generation entirely on the user's machine. No hosted language model is used at runtime.
all-MiniLM-L6-v2 (Sentence-Transformers)	The embedding model used for semantic search and for placing papers on the literature map. It runs locally and is bundled with the application.

Date: 17 August 2026

ACKNOWLEDGEMENTS

We thank our guide, Dr. New Begin, for supervision throughout the project, and the Department of Computer Science, CHRIST (Deemed to be University), for the facilities and the opportunity to undertake this work.

We also thank the twelve testers who installed the pre-release builds on their own machines. Several defects recorded in this report, including a graphics configuration that reported no accelerator on hardware that had one, were found only because the application was run somewhere other than the machines it was written on.

ABSTRACT

A literature review requires a researcher to collect papers, read them, find what they collectively argue, identify what they do not yet answer, and write the result. The tools that assist with this are almost without exception hosted services: the document is uploaded, a remote language model reads it, and the answer is returned. That arrangement is unavailable to a researcher whose material is unpublished, institutionally restricted or personally identifiable, and it carries a per-query cost that a student project cannot meet.

ThinkStack performs the same work with no network. A quantised small language model runs on the user’s own processor; embeddings are computed locally; the typesetting engine is bundled with the installer. The application ingests PDFs, extracts their bibliographic metadata from page geometry, indexes them for semantic search, places them on a map by meaning, identifies gaps between them, and provides a $\text{\LaTeX}$ editor that compiles offline.

The system is a desktop application built on a Python FastAPI backend, a React interface and a Rust (Tauri) shell, distributed as signed installers for Linux, Windows and macOS. It has been released publicly through fifteen versions.

Two results are reported. First, metadata extraction was rebuilt to read positioned glyphs rather than flattened text, and measured against 56 papers sampled at random across 15 arXiv subject classifications, using the archive’s own catalogue as ground truth: **92.9% of titles and 85.7% of author lists exactly correct**, against a prior implementation that produced no usable author list on any published paper tested. Second, the constraint the project was built under, competitive quality on a consumer laptop with no network and no accelerator, is itself the contribution, since the established tools in this area assume a server.

Keywords: small language models, edge inference, offline systems, literature review, retrieval-augmented generation, document layout analysis, $\text{\LaTeX}$.

Contents

List of Abbreviations	v
1 INTRODUCTION	1
1.1 Project Description	1
1.2 Existing System	1
1.2.1 Hosted research assistants	1
1.2.2 Reference managers	1
1.2.3 Hosted L ^A T _E X editors	2
1.2.4 What none of them do	2
1.3 Objectives	2
1.4 Purpose, Scope and Applicability	2
1.4.1 Purpose	2
1.4.2 Scope	2
1.4.3 Applicability	3
1.5 Overview of the Report	3
2 SYSTEM ANALYSIS	4
2.1 Problem Definition	4
2.2 Requirements Specification	4
2.2.1 Functional Requirements	4
2.2.2 Non-Functional Requirements	5
2.3 Block Diagram	7
2.4 System Requirements	7
2.4.1 User Characteristics	7
2.4.2 Software and Hardware Requirements	7
2.4.3 Constraints	8
2.5 Conceptual Models	8
2.5.1 Data Flow Diagram	8
2.5.2 Entity–Relationship Diagram	9
3 SYSTEM DESIGN	10
3.1 System Architecture	10
3.2 Module Design	10
3.3 Database Design	11
3.3.1 Tables and Relationships	11
3.3.2 Data Integrity and Constraints	12
3.4 System Configuration	12
3.5 Interface Design and Procedural Design	12
3.5.1 User Interface Design	12
3.5.2 Application Flow and Class Diagram	13
3.6 Reports Design	14

4	IMPLEMENTATION	15
4.1	Implementation Approaches	15
4.1.1	Measure, then move	15
4.1.2	Repair rather than refuse	15
4.1.3	Record the reasoning, not the change	15
4.2	Coding Standard	15
4.3	Key Implementations	16
4.3.1	Metadata extraction from page geometry	16
4.3.2	Sizing the work to the machine	17
4.3.3	The guard that tested for silence	17
4.3.4	A project is a directory	17
4.3.5	Citations, and where the authority for a key lives	17
4.3.6	Correcting what extraction gets wrong	18
4.3.7	What other machines found	18
4.4	Screenshots	19
5	TESTING	22
5.1	Testing Approaches	22
5.2	Test Cases	22
5.3	Test Reports	24
5.3.1	Metadata extraction, measured on unseen papers	24
5.3.2	Suite results	24
6	CONCLUSION	25
6.1	Design and Implementation Issues	25
6.1.1	A guard cannot test for the absence of a failure it never produces	25
6.1.2	Constants standing in for measurements	25
6.1.3	Measuring on material one has chosen	25
6.1.4	Replacing a running interface costs double	25
6.2	Advantages and Limitations	25
6.2.1	Advantages	25
6.2.2	Limitations	26
6.3	Future Scope	26
	References	28

List of Tables

2.1	Functional requirements	4
2.2	Non-functional requirements	6
2.3	Hardware and software requirements	7
3.1	Architectural layers	10
3.2	Module decomposition	11
3.3	Persistent stores	11
5.1	Test layers	22
5.2	Representative test cases	22
5.3	Extraction accuracy, curated versus randomly sampled papers	24
5.4	Suite results at the time of writing	24

List of Figures

2.1	Block diagram. The shell supervises the backend process and hosts the interface; every arrow is a local call.	7
2.2	Level-1 data flow. No process crosses the machine boundary.	8
2.3	Entities and their relationships. A document owns its chunks; an analysis run covers a selection of documents and produces gaps.	9
3.1	Application flow: one paper, from drop to citable.	13
3.2	The ingestion data model. TextSpan carries the baseline, which is what makes a small-capitals title recoverable.	14
4.1	Library. The folio above the rule counts the collection; the three panels report which model serves which task, which papers are missing a field, and which drafts are open. The shelf lists the papers themselves.	19
4.2	LitGraph. Papers are placed by semantic similarity, themes are drawn as regions, and links show the strongest relations. The two clusters are the change-point detection and post-quantum cryptography literatures; the space between them is where the gap finder looks. Search, gap finding and the encrypted vault are reached from this one view.	20
4.3	Scribe. The project is a directory: the tree on the left, the source in the middle, and the compiled PDF on the right, rebuilt a moment after typing stops.	20
4.4	Citing from the editor. Typing the word cite opens the library beneath the cursor, with the key each paper would be cited by shown at the right and papers already in this project's bibliography marked. Enter replaces the word with the citation command.	21
4.5	Bench. What the machine can run, measured rather than assumed, including a graphics device that is present but not in use because the shipped build is processor-only.	21

LIST OF ABBREVIATIONS

Abbreviation	Expansion
ADR	Architecture Decision Record
API	Application Programming Interface
CI	Continuous Integration
CPU	Central Processing Unit
CRF	Conditional Random Field
DFD	Data Flow Diagram
DOI	Digital Object Identifier
ER	Entity-Relationship
GGUF	GPT-Generated Unified Format (quantised model container)
GPU	Graphics Processing Unit
GUI	Graphical User Interface
JSON	JavaScript Object Notation
LLM	Large Language Model
NFC/NFKC	Unicode Normalisation Forms (Composition / Compatibility)
OCR	Optical Character Recognition
PDF	Portable Document Format
RAG	Retrieval-Augmented Generation
REST	Representational State Transfer
SLM	Small Language Model
SRS	Software Requirements Specification
TTS	Text-to-Speech
UI/UX	User Interface / User Experience
VRAM	Video Random Access Memory

Chapter 1

INTRODUCTION

1.1 Project Description

ThinkStack is a desktop research assistant that reads a collection of academic papers and helps its user review and write about them, running entirely on the machine it is installed on.

The user adds PDFs. Each is parsed, its bibliographic metadata extracted, its text divided into passages, and each passage converted into a vector by an embedding model. From that index the application provides four capabilities, which the interface presents as four screens:

- **Library**, the collection: what has been ingested, what has been read, what is incomplete, and what the machine is running.
- **LitGraph**, the collection drawn as a map, papers placed by meaning rather than by folder, with themes as territories and gaps as markers. Search, analysis and gap-finding operate on a selection made on the map itself.
- **Scribe**, a $\text{\LaTeX}$ editor with a bundled typesetting engine, an auto-healing compiler, and a natural-language-to-markup path that is grounded on the user's own papers.
- **Bench**: what the machine can run, which model serves which task, and how to add or remove models.

Every inference is local. The single optional network request is a version check against the releases page, which transmits no user data and fails silently when offline.

1.2 Existing System

Tools that assist with literature review fall into three groups, and each is unusable under at least one of this project's constraints.

1.2.1 Hosted research assistants

Elicit, SciSpace, Consensus and comparable services upload the document and answer from a remote large language model. They are capable and, for a large class of users, the correct choice. They are unavailable to a researcher whose material cannot leave the machine: unpublished results, institutionally restricted corpora, or documents containing personal data, and they charge per query, which a student project cannot sustain.

1.2.2 Reference managers

Zotero and Mendeley manage citations well and do not read the papers. They organise; they do not answer questions about content, find gaps, or write.

1.2.3 Hosted L^AT_EX editors

Overleaf compiles in the cloud, which introduces a queue, a time limit and a network dependency, and requires an account. It also has no knowledge of the user’s library, so a citation must be fetched by hand from elsewhere.

1.2.4 What none of them do

None runs the model on the user’s own machine. The gap ThinkStack addresses is not a missing feature in any one of these tools; it is that the entire category assumes a server, and therefore assumes the document may be sent to one.

1.3 Objectives

1. Run every inference, embedding, summarisation, claim extraction, gap analysis and text generation, on the user’s own hardware, with no hosted language model at any point.
2. Remain fully functional with no network connection.
3. Extract bibliographic metadata accurately enough that it can label a paper elsewhere in the application and populate a citation.
4. Represent a collection as a structure rather than a list, so that the relationships between papers and the gaps between them are visible.
5. Compile L^AT_EX locally, with a bundled engine, so that writing has no queue, no account and no time limit.
6. Size the workload to the machine, so that the application is usable on a consumer laptop rather than requiring an accelerator.
7. Distribute as installers for Linux, Windows and macOS, verified by running the packaged application rather than by a green build.

1.4 Purpose, Scope and Applicability

1.4.1 Purpose

To make the assistance that hosted research tools provide available to researchers who cannot use them, by moving the computation to the user’s device rather than the user’s documents to a provider’s.

1.4.2 Scope

The system covers ingestion of text-bearing PDFs, metadata extraction, semantic search, summarisation, claim extraction, thematic clustering, gap analysis, a literature map, at-rest encryption of stored documents, local L^AT_EX compilation, and model management including acquisition from Hugging Face.

It does not cover scanned documents with no text layer, which require OCR that the application does not ship; collaborative editing; or synchronisation between devices, which would contradict the premise.

1.4.3 Applicability

The system applies wherever a literature review must be conducted on material that cannot be uploaded, or where per-query cost or network availability rules out a hosted tool: postgraduate and doctoral research, institutions with data governance requirements, and students without a budget for inference.

1.5 Overview of the Report

Chapter 2 states the problem formally and specifies the functional and non-functional requirements, the system requirements, and the conceptual models. Chapter 3 presents the architecture, the module decomposition, the data design and the interface design. Chapter 4 describes the implementation, the coding standards adopted, and the parts of the system whose implementation is worth recording in detail. Chapter 5 reports the testing strategy, the test cases and their results. Chapter 6 concludes with the design and implementation issues encountered, the advantages and limitations of the result, and the future scope of the work.

Chapter 2

SYSTEM ANALYSIS

2.1 Problem Definition

Given a collection of research papers held locally, the system must let a researcher answer four questions without any document leaving the machine:

1. **What do I have?:** identify each paper accurately enough to name it, and report what the application has read of it.
2. **What does it say about X ?:** retrieve the passages that answer a question posed in natural language.
3. **What does the collection not answer?:** identify claims that are asserted but unsupported, and questions the corpus raises but does not address.
4. **How do I write it up?:** compose and typeset the result, citing the collection.

Each is constrained by the hardware the work must run on. A consumer laptop with 8–16 GB of memory and no usable accelerator is the target, not the degraded case.

2.2 Requirements Specification

2.2.1 Functional Requirements

Table 2.1: Functional requirements

ID	Requirement	Description
FR-1	Ingest a PDF	Accept one or many PDFs by drag-and-drop or file picker; extract text; report failure per file.
FR-2	Extract metadata	Determine title, authors, year and abstract from the document, without a network lookup.
FR-3	Correct metadata	Allow a wrong title, author list or year to be corrected in place, and propagate the correction to every stored passage.
FR-4	Chunk and embed	Divide text into passages and compute an embedding for each, locally.
FR-5	Semantic search	Return ranked passages for a natural-language query, grouped by paper.
FR-6	Summarise	Produce a summary of one or more selected papers.
FR-7	Extract claims	Produce the assertions a paper makes, as structured output.

ID	Requirement	Description
FR-8	Cluster themes	Group the collection into themes derived from the embeddings.
FR-9	Find gaps	Identify unsupported claims and unaddressed questions across a selection, and retain past runs.
FR-10	Draw the map	Place papers in two dimensions by semantic similarity; show themes as regions and gaps as markers.
FR-11	Read a paper	Display the original PDF with matched passages marked.
FR-12	Encrypt at rest	Encrypt a stored document under a user password and decrypt on demand.
FR-13	Edit and compile	Provide a L ^A T _E X editor and compile to PDF locally, recovering from common errors.
FR-14	Manage project files	Treat a paper as a directory; support upload, rename, move, copy and delete within it.
FR-15	Generate markup	Convert a natural-language instruction into L ^A T _E X, optionally grounded on selected papers.
FR-16	Manage models	List, import, assign, download and remove models; route each task to a model.
FR-17	Report capability	Report what the machine can run, measured rather than assumed.
FR-18	Update in place	Check for a new version and install it on request.
FR-19	Report progress	Show what background analysis is doing and how far through it is.
FR-20	Stop ingestion	Interrupt a running batch without leaving a document partly ingested.
FR-21	Cite a paper	Insert a citation to a paper in the library from within the editor, and maintain the project's bibliography file.
FR-22	Resolve citations	Compile citations to numbered references and generate the reference list, without the author declaring a bibliography.

2.2.2 Non-Functional Requirements

Table 2.2: Non-functional requirements

ID	Attribute	Requirement
NFR-1	Privacy	No document, prompt or embedding is transmitted. The only outbound request is a version check carrying no user data.
NFR-2	Offline operation	Every feature except the update check functions with no network.
NFR-3	Memory	The model is sized against <i>available</i> memory, not installed memory, so the process is not terminated mid-analysis.
NFR-4	Responsiveness	Ingestion returns before analysis completes; the interface reports the queue rather than blocking.
NFR-5	Portability	One codebase produces installers for Linux, Windows and macOS.
NFR-6	Recoverability	A malformed figure or a missing package must not cost the author the compiled document.
NFR-7	Integrity	Encryption is authenticated, so tampering is detected at decryption rather than producing plausible but corrupted output.
NFR-8	Maintainability	Every architectural decision is recorded with its reasoning; dependencies are pinned exactly.
NFR-9	Verifiability	A release is validated by running the packaged application, not by a passing build.
NFR-10	Accessibility	Text meets contrast requirements against its substrate; controls are reachable by keyboard.

2.3 Block Diagram




Figure 2.1: Block diagram. The shell supervises the backend process and hosts the interface; every arrow is a local call.

2.4 System Requirements

2.4.1 User Characteristics

The intended user is a researcher, postgraduate student or academic who reads and writes papers. Familiarity with $\text{\LaTeX}$ is assumed for Scribe but not for the rest of the application. No knowledge of language models, embeddings or model quantisation is assumed, which is why Bench reports what the machine can run in plain terms and the application ships with a working model rather than asking the user to choose one.

2.4.2 Software and Hardware Requirements

Table 2.3: Hardware and software requirements

	Minimum	Recommended
Operating system	Windows 10, macOS 12, or a Linux distribution with GTK/WebKit	Any of the above, current
Processor	64-bit dual core	Quad core or better
Memory	8 GB	16 GB
Disk	2 GB free	5 GB free
Accelerator	None required	Vulkan-capable GPU, optional

	Minimum	Recommended
Network	Not required	Required only to download an additional model

The development and evaluation machine was a 16 GB laptop with no usable accelerator, running Fedora. Every figure reported in Chapter 5 was measured on it.

2.4.3 Constraints

- **No hosted inference.** This forecloses the largest models and is the constraint from which most of the design follows.
- **Installer size.** A single release asset must remain under GitHub’s 2 GiB limit, which is why one small model is bundled and a larger one is fetched on consent.
- **A shared, serialised model.** One model is resident at a time, so concurrent analyses queue rather than compete for memory.
- **No OCR.** A scanned PDF with no text layer yields nothing, and no later stage can recover from that.

2.5 Conceptual Models

2.5.1 Data Flow Diagram




Figure 2.2: Level-1 data flow. No process crosses the machine boundary.

2.5.2 Entity–Relationship Diagram




Figure 2.3: Entities and their relationships. A document owns its chunks; an analysis run covers a selection of documents and produces gaps.

Chapter 3

SYSTEM DESIGN

3.1 System Architecture

The application is four layers and a shell, and the direction of dependency is strictly downward: the interface knows the API, the API knows the domain, the domain knows infrastructure, and nothing knows the layer above it.

Table 3.1: Architectural layers

Layer	Responsibility
frontend/	The interface. React 19 and Vite, one theme, four screens. Holds no business rule: anything derived from global state is read from the backend rather than recomputed.
api/	REST endpoints. Validates, orchestrates, and returns; contains no domain logic of its own.
domain/	The capabilities, one package each: <code>ingestion</code> , <code>knowledge_base</code> , <code>search</code> , <code>analysis</code> , <code>litgraph</code> , <code>paper_writer</code> , <code>model_manager</code> . Pure Python, no framework imports.
infrastructure/	Everything that touches the machine: the inference client, the vector store, the capability model, file management, the job queue, encryption.
src-tauri/	The desktop shell, in Rust. Starts the backend as a supervised child process, waits behind a loading screen, and stops it when the window closes.

Decision recorded. One module answers “what can this machine do”. Before `infrastructure/capability.py` existed, that question was answered in about ten places and two of them disagreed, the Rust diagnosis chose GPU layers and a Python module computed them again, with the Rust answer winning at runtime. The Python one was therefore unreachable code that still looked authoritative. The split is now: *Rust detects, Python decides*.

3.2 Module Design

Table 3.2: Module decomposition

Module	Contents
ingestion	<code>pdf_parser</code> (text and positioned spans), <code>layout</code> (geometry: rows, cells, normalisation), <code>layout_metadata</code> (title and authors from geometry), <code>metadata_extractor</code> (orchestration and fallbacks), <code>chunker</code> .
knowledge_base	<code>embedding_service</code> , <code>repository</code> (storage and retrieval of chunks with their metadata).
search	Semantic search over chunk embeddings, with optional grouping by document.
analysis	<code>summarizer</code> , <code>claim_extractor</code> , <code>theme_clusterer</code> , <code>precompute</code> (queues work at ingest so the canvas is ready before it is opened).
litgraph	<code>graph_builder</code> : positions, themes and gap markers derived on request from embeddings and past runs.
paper_writer	<code>compiler</code> (Tectonic or pdflatex, with error recovery), <code>files</code> (a project is a directory, with a path boundary), <code>indexing</code> .
model_manager	<code>registry</code> , <code>catalog</code> , <code>router</code> , <code>downloader</code> , <code>manifest</code> .

3.3 Database Design

3.3.1 Tables and Relationships

The system stores five things, all under the user’s data directory. There is no relational database: the store is a vector index with per-record metadata, plus JSON documents and directories on disk. That is a deliberate choice, a relational engine would add an installed dependency to an application whose premise is that it runs on a laptop with nothing configured, and the consequence is that referential integrity is enforced in the repository layer rather than by a foreign key.

Table 3.3: Persistent stores

Store	Form	Contents
Vector store	Embedded chunks	<code>chunk_id</code> , <code>doc_id</code> , <code>text</code> , <code>embedding</code> , <code>page_number</code> , <code>chunk_index</code> , <code>token_count</code> , and the paper’s <code>title</code> , <code>authors</code> , <code>year</code> .
Papers directory	Original PDFs	The uploaded file, unmodified, addressed by <code>doc_id</code> .
Analysis cache	JSON	Per document: summary and extracted claims.

Store	Form	Contents
Run histories	JSON	Analysis and gap runs, newest first, capped.
Paper projects	Directories	One directory per paper, containing <code>main.tex</code> , figures and build artefacts.

3.3.2 Data Integrity and Constraints

- Bibliographic metadata is duplicated onto every chunk so that a search result can name its paper without a second lookup. The cost is that correcting a title is not one write but a write to every chunk of the document; missing one would let the same paper answer to two titles depending on which passage matched.
- Deleting a document deletes its chunks, its stored PDF and its cached analysis together.
- Encryption is applied to chunk text in the store. The original PDF is *not* encrypted, so the route that serves it refuses to do so for an encrypted document: serving it would hand back exactly what the encryption exists to withhold.
- A gap run supersedes the one before it: the newest run is the library's current state, not the sum of all runs.

3.4 System Configuration

Configuration is derived rather than declared wherever possible. The capability model measures the machine at startup and every downstream number: context size, GPU layers, output token budget, how much of a paper fits in one prompt is computed from that measurement rather than set in a file. Dependencies are pinned exactly, and a check in CI fails the build if a range appears, since a range lets the same commit build a different program tomorrow.

3.5 Interface Design and Procedural Design

3.5.1 User Interface Design

The interface is one theme, described internally as *Paper and Ink*: a sheet of paper on a desk, worked in ink. Four surfaces, one ink ramp of four steps, a single 2px radius and two shadows. Type is a six-step scale on a fluid root, so every measurement grows with the window rather than sitting at a fixed size on a large display.

Three principles govern it, and each was arrived at by correcting its opposite:

1. **The interface reads the backend's answer; it never recomputes it.** A panel that decided for itself which model was in use displayed the wrong one for every release in which it existed.
2. **A figure earns its place by answering a question the reader would ask.** Counts that described the implementation, bytes on disk, the number of index fragments, were removed rather than restyled.

3. **Nothing is withheld.** A collapsed navigation panel keeps a rail rather than disappearing, because at zero width the controls it still owed the reader had to be drawn on top of the page.

3.5.2 Application Flow and Class Diagram

Figure 3.1 traces one paper from the moment it is dropped on the window to the moment it can be cited, which is the path that crosses every module in the system.




Figure 3.1: Application flow: one paper, from drop to citable.

Steps 1–8 complete before the upload request returns. Step 9 is deliberately *not* awaited: it costs tens of seconds per paper, and awaiting it once put a fifty-second model call inside an HTTP request, so uploading three papers appeared to hang the application for three minutes. The work still happens at ingest time, so the map is ready before it is opened, it simply happens after the response, with the queue reported in the interface.




Figure 3.2: The ingestion data model. **TextSpan** carries the baseline, which is what makes a small-capitals title recoverable.

3.6 Reports Design

The application produces three outputs a user keeps: the compiled PDF from Scribe; the gap-analysis run, retained in history and reopenable; and the literature map, which is regenerated from the current index on every visit rather than stored, so there is no stale graph to invalidate.

Chapter 4

IMPLEMENTATION

4.1 Implementation Approaches

Three approaches shaped the implementation and are worth stating because each replaced a plausible alternative.

4.1.1 Measure, then move

Work was moved between languages only for a measured reason. Startup diagnosis was rewritten in Rust because the Python implementation imported a deep-learning framework purely to discover whether a GPU existed, costing seconds on every launch. Model discovery was measured at approximately 9ms and left in Python, because the same argument does not apply to code that is already fast.

The same rule prevented a larger mistake. When metadata extraction was found to be wrong, rewriting the pipeline in a systems language was considered. Profiling one paper gave: extraction 6.0ms, metadata 1.1ms, chunking 0.2ms, embedding and search 3073.8ms. Everything a rewrite would have touched was 0.3% of runtime, and the remainder was already compiled native code. The defect was the question being asked, not the language asking it.

4.1.2 Repair rather than refuse

A small language model produces markup that is frequently almost correct. The compiler therefore inserts missing package declarations, wraps bare fragments into complete documents, and progressively neutralises a failing environment and retries, so one malformed figure does not cost the author the document.

4.1.3 Record the reasoning, not the change

Every architectural decision is recorded in `docs/ADR.md` with its context and its consequences. The log reached 40 entries. Its value was not apparent while it was being written and became apparent during the dependency audit, when the reasoning behind earlier choices could be read rather than reconstructed.

4.2 Coding Standard

- **Python:** `ruff` in CI, no warnings tolerated; type hints on module boundaries; no framework imports inside `domain/`.
- **JavaScript:** ESLint with the React hooks rules; zero errors required to merge. The hooks linter caught a genuine defect during this work: a ref read during render.
- **Comments,** a comment states the constraint a maintainer must not break, not what the line does. Rationale belongs in the decision log. A prose-reduction pass took

comment and docstring density from 32% to 27% of source lines, deleting restatements of type hints and of decisions recorded elsewhere.

- **Commits**, the reasoning goes in the commit message, and the history is treated as part of the deliverable.
- **Branching**: `main` (released), `beta` (testers), `dev` (integration), and short-lived `feat/` / `fix/` branches. Nothing reaches `main` without passing through the other two.

4.3 Key Implementations

4.3.1 Metadata extraction from page geometry

This is the implementation the project is most defined by, and it is the one that was rebuilt.

The original implementation read the document as a flat string and applied regular expressions: the title was any of the first ten lines longer than ten characters; an author was any line beginning with two capitalised words; the year was the first four-digit number. Measured against published papers with known metadata, it produced no usable author list on any of them. On *Attention Is All You Need* it returned the publisher’s copyright notice as the title, listed the paper’s own title and its authors’ employers as authors, and reported the year as 2014: three years before the paper existed.

The cause is that a PDF is not text. It is positioned glyphs, each carrying a size, a bounding box and a writing direction, and `extract_text()` discards all of it. Typesetting convention places the title in the largest horizontal text at the top of page 1, with the authors on the rows beneath, a signal that is present in the file and thrown away before extraction begins.

The rebuilt implementation is two modules. `layout.py` knows about typesetting and nothing about papers: it groups glyphs into rows and divides each row into cells. `layout_metadata.py` decides what a title or an author is, as a set of named rules.

Splitting rows into cells is the single highest-value step. The PDF library reports a four-column author row as one line, so read as text “Kaiming He Xiangyu Zhang Shaoqing Ren Jian Sun” is an eight-word phrase and no filter can recover four names from it. Read as positions it is four cells. Ablated, that one rule is worth 7 of 7 papers against 1 of 7.

Three further findings were geometric rather than bibliographic:

- Rows must be grouped by *baseline*, not by bounding-box top. A small-capitals title is one line set in two sizes; grouped by the top, *ADAM, A Method for Stochastic Optimization* extracted as “A, A M S O”.
- A font change mid-word produces two spans with no gap between them. Inserting a space there yields “A DAM”.
- T_EX emits an accent as a separate glyph *before* its letter, so “Dollár” arrives as three pieces and prints as “Doll ár”, leaving the surname unsearchable.

Affiliations are removed without naming employers, which would be an endless and dating list. Three rules do it instead: a structural word (*university, institute, research*) disqualifies an entire row; a short all-capitals token beside ordinary words is an acronym; and a phrase printed more than once in the author region is an address, because a shared affiliation is printed once per author while a name is printed once per paper. The third

requires no vocabulary at all and is the only rule that can separate “United Kingdom” from “Kaiming He”, which are indistinguishable by spelling.

4.3.2 Sizing the work to the machine

`capability.py` answers what the machine can do, and every derived number comes from it. Two defects were fixed by centralising it:

- **Summarisation could not fit in its own context window.** 6000 characters of paper plus a 1024-token reply needs roughly 2650 tokens; a low-capability machine is given 2048. The request failed before the model was asked to think, and the error handler then told the reader the response “could not be read”, which was untrue.
- **Every Apple Silicon machine was pinned to the CPU.** GPU layers were decided by `has_cuda && vram_gb >= 2.0`; Apple Silicon reports neither.

4.3.3 The guard that tested for silence

A recurring defect class is worth recording as an implementation lesson. Three separate failures had the same shape: *a check that tests for silence, when the failure mode is confident wrongness.*

A small-language-model fallback existed for papers the regular expressions could not handle, guarded on the extracted title being *empty*. The regular expressions never returned empty, they returned wrong, so the guard never fired on a single paper it existed to rescue. Similarly, graphics advice was gated on a VRAM figure that only one vendor’s tool reports, so machines with a working accelerator were told they had none; and a release marked `draft: false` was published as a draft, with every job reporting green.

The correction in each case is the same: test for the shape of a right answer, not for the absence of one.

4.3.4 A project is a directory

Compilation always ran with the project directory as its working directory, so relative paths in `\includegraphics` resolved correctly, but there was no way to put a second file into that directory, so a figure could be referenced and never supplied. Exposing the directory required a path boundary: this API is reachable from the webview, which the shell treats as remote content, so a filename arriving at it is untrusted input. Every operation resolves its argument against the project directory and refuses anything landing outside. Resolution happens *before* the check rather than a string test for `..`, because that is what catches a symbolic link.

4.3.5 Citations, and where the authority for a key lives

The editor could compile a document but could not cite one. The interaction chosen puts the library where the author already is: typing the word `cite` opens a list beneath the cursor, further typing filters it on title, author, year or key, and Enter replaces the word with a citation command while writing the corresponding entry into the project’s bibliography file. Selecting nothing leaves the word in place, and an ordinary word compiles.

Citation numbers are not tracked, and the omission is deliberate. BibTeX rennumbers every citation on each recompile and orders the reference list according to the chosen style, so emitting keys obtains that behaviour for nothing, whereas tracking numbers would maintain a second and less reliable copy of it.

The design question that remained was which side holds the authority for a key. Keys are derived from metadata, so a collision between two papers is resolved differently depending on what else the library contains, and a key derived again later can differ from the one already written into the author's source. The bibliography file is therefore authoritative: keys already present in it are honoured and never recomputed, and only documents the file has not seen are assigned new ones. Each entry carries an identifier field that bibliography styles ignore, since a style reads only the fields it declares, which preserves the link back to the library record even when the file has been edited by hand.

Two defects underneath had to be corrected first, and both printed rather than raised. Author lists were stored by joining them with commas, but in a bibliography the comma separates a surname from a given name, so eight authors were read back as sixteen fragments. And a document containing citations but no bibliography command produced no references and no error, because the command that prints the reference list is also the one that causes the bibliography processor to run at all.

4.3.6 Correcting what extraction gets wrong

Extraction identifies titles correctly for about 93 percent of papers and author lists for about 86 percent, so roughly one paper in seven carries a wrong record. That record names the paper on the map, in every search result and in every reference generated from it, which means a citation feature built above it inherits the error.

Correction belongs to the library rather than to the editor. The library holds the record, so a correction there applies to every subsequent citation, every node on the map and every search result at once, while a correction made inside one document's bibliography applies to that document alone. The editor needs no separate facility, because a project's bibliography is already a file in its own tree and the identifier field described above lets it be edited by hand without breaking the link back to the record. Placing the same field in two editable surfaces would create two sources of truth for one value.

4.3.7 What other machines found

Three defects reached the application through people rather than through the suites, and each is the same lesson in a different place: an automated test observes the layer it was written for, and every one of these lived somewhere else.

The reader failed on Windows because the server described a JavaScript module as plain text. Python builds part of its media type table from the Windows registry, which carries no entry for the extension in question; a browser then refuses to execute the file it has just successfully downloaded. No test in the suite reads a response header, and on the two platforms used in development the table is correct, so the defect was invisible in every sense available to us.

The gap finder listed the same finding more than once, because each entry the model returned was given a fresh identifier before anything compared them. The model repeats itself, as small models do; treating its reply as a set of findings rather than as a list of rows removes the duplicate without discarding the evidence the duplicate carried.

Selecting a point on the map required more precision than a pointer reliably has. The click target was the drawn point, which is about six pixels once the map is scaled to show a whole collection. The drawing was deliberately left alone and an invisible target now carries the selection, held at a constant size on screen as the view zooms. Enlarging the drawing, or opening the map already magnified, were both considered and rejected: the first turns the map into a bubble chart and the second hides the structure the map exists to show.

4.4 Screenshots

The following are captured from the running application on the development machine, against a library of 20 papers.

The screenshot shows the ThinkStack application interface. The top bar displays statistics: 21 papers, 20 Read, 4 gaps, 464 chunks, and a Refresh button. The sidebar on the left contains navigation options: Library (selected), LitGraph, Scribe, and Bench. The main content area is divided into three panels: MODELS IN USE, NEEDS ATTENTION, and BEING WRITTEN. Below these panels is a search filter and a list of papers.

MODELS IN USE	NEEDS ATTENTION	BEING WRITTEN
General	Evidence for long-tailed distributions in the Inte...	no year citation test pdf
Analysis	The impact of data-heavy, post-quantum TLS 1...	no year bundle-validation pdf
Gap finding	Measuring Latency Variation in the Internet	no year untitled pdf
LaTeX writing		uxtest pdf

Runs on this machine. [Change in Bench](#) 3 of 21. [Show only these](#) 23 in progress. [Continue writing](#)

Drop PDFs here, or [browse](#)

Q. Filter 21 papers by title, author or filename

Title	Author	Status	Year
State-of-the-Art in Sequential Change-Point Detection	Aleksey S. Polunchenko, Alexan...	read	2011
Large-Scale Characterization and Segmentation of Internet Path Delays ...	Maxime Mouchet, Sandrine Vaton...	read	2019
Evidence for long-tailed distributions in the Internet	Allen B. Downey	read	-
Lagged Exact Bayesian Online Changepoint Detection with Parameter Es...	Michael Byrd, Linh Nghiem, Jin...	read	2017

Figure 4.1: Library. The folio above the rule counts the collection; the three panels report which model serves which task, which papers are missing a field, and which drafts are open. The shelf lists the papers themselves.




Figure 4.2: LitGraph. Papers are placed by semantic similarity, themes are drawn as regions, and links show the strongest relations. The two clusters are the change-point detection and post-quantum cryptography literatures; the space between them is where the gap finder looks. Search, gap finding and the encrypted vault are reached from this one view.




Figure 4.3: Scribe. The project is a directory: the tree on the left, the source in the middle, and the compiled PDF on the right, rebuilt a moment after typing stops.




Figure 4.4: Citing from the editor. Typing the word `cite` opens the library beneath the cursor, with the key each paper would be cited by shown at the right and papers already in this project’s bibliography marked. Enter replaces the word with the citation command.




Figure 4.5: Bench. What the machine can run, measured rather than assumed, including a graphics device that is present but not in use because the shipped build is processor-only.

Chapter 5

TESTING

5.1 Testing Approaches

Testing is in four layers, and the reason for the fourth is the interesting one.

Table 5.1: Test layers

Layer	Count	What it establishes
Backend unit tests	984	Domain and infrastructure behaviour, run on every push and in CI.
Interface unit tests	198	Component and utility behaviour under jsdom.
Contract tests	5	That every route the interface calls is a route the backend defines: both sides parsed from source.
Packaged-application validation	:	That the <i>installed</i> application starts and works, run against the built installer on each platform.

The fourth layer exists because of a specific failure. Version 1.0.0 passed every test and every build, was published, and did not start: the frozen backend silently omitted the inference library’s shared objects and the embedding model’s data, because neither package ships a hook for the freezer. A test suite that runs from a source checkout cannot see that. `validate_bundle.py` now runs against the packaged application and checks outcomes rather than statuses: for example, inspecting a compiled PDF for an image object rather than trusting a “compiled” result, because a compile that silently drops the figure also reports success.

5.2 Test Cases

Representative cases from each area; the full suite is in `tests/`.

Table 5.2: Representative test cases

ID	Case	Expected	Result
T-01	Title is the largest horizontal text at the top of page 1	Exact title returned	Pass
T-02	Rotated archive stamp larger than the title	Stamp ignored	Pass
T-03	Small-capitals title (two sizes, one baseline)	Whole title, no split	Pass
T-04	Accent stored before its letter	Surname recomposed	Pass
T-05	Affiliation row naming two employers, only one listed	Whole row rejected	Pass

ID	Case	Expected	Result
T-06	Phrase repeated in the author region	Treated as an address	Pass
T-07	Year from an archive identifier vs. a stray four-digit number	Identifier wins	Pass
T-08	Paper stating no year	Year left empty	Pass
T-09	Plausible title, no authors	Model fallback invoked	Pass
T-10	Model returns an empty field	Existing good value retained	Pass
T-11	Correction propagates to every chunk	All chunks updated	Pass
T-12	Correction with a blank title	Rejected, 400	Pass
T-13	Correction to an absent document	404, nothing written	Pass
T-14	Ingest continues when the analysis queue fails	Document retained	Pass
T-15	PDF route for an encrypted document	Refused, 403	Pass
T-16	Filename escaping the project directory	Refused	Pass
T-17	Sidebar collapse does not swallow the click that caused it	Click delivered	Pass
T-18	Update prompt where the platform cannot ask	Not installed	Pass
T-19	Every interface call resolves to a defined route	No orphans	Pass
T-20	Version never decreases	Lower version refused	Pass
T-21	An author list survives storage	Eight names read back as eight	Pass
T-22	A key already in the bibliography is not recomputed	Same key returned	Pass
T-23	Citing the same paper twice	One entry, same key	Pass
T-24	An entry written by hand survives a citation	Retained	Pass
T-25	Listing what can be cited writes nothing	No file created	Pass
T-26	A citing document without a bibliography	Command supplied, references typeset	Pass
T-27	Special characters in a title reach the bibliography	Escaped	Pass
T-28	A JavaScript module is served as JavaScript	Never <code>text/plain</code>	Pass
T-29	The same gap returned twice	One finding, evidence merged	Pass
T-30	A selectable target on the map	28 px at every usable zoom	Pass

5.3 Test Reports

5.3.1 Metadata extraction, measured on unseen papers

The development set was 18 papers chosen by the authors, on which the rebuilt extractor was exactly correct on every title, author list and year. That figure was then discarded as unsafe, because a measurement taken over material one has selected oneself describes the selection rather than the method.

A second measurement was constructed in which neither the papers nor the expected answers were chosen by the authors: 56 papers sampled across 15 arXiv subject classifications through the public interface, scored against the archive’s own catalogue records.

Table 5.3: Extraction accuracy, curated versus randomly sampled papers

Field	Curated (18)	Random (56)
Title exactly correct	100%	92.9%
Author list exactly correct	100%	85.7%
Year exactly correct	100%	92.9%
All three simultaneously	100%	76.8%

The first random measurement was 62.5% on all three fields. The entire difference from the final figure was typesetting conventions that had not been examined: one physics template spaces a centred author line widely enough to be indistinguishable from column separation; another runs from the affiliations into the abstract with no heading, so the search for authors continued into the body; and a margin threshold introduced to discard an archive stamp was discarding the first line of any centred title wide enough to begin left of it.

Of the 13 remaining failures, three are not extraction errors: the page reads “M. GRENDÁR” and “Audrey Quessadda-Vial” where the archive record says otherwise.

5.3.2 Suite results

Table 5.4: Suite results at the time of writing

Suite	Tests	Passing	Runtime
Backend (pytest)	1 085	1 084	~14 s
Interface (vittest)	222	222	~5 s
Contract	5	5	<1 s
Route reachability	16	16	<2 s

Route reachability is distinct from the contract test and was added when the interface was replaced. Parsing both sides as text proves the paths *match*; calling them against a running backend proves the handlers *run*.

Chapter 6

CONCLUSION

6.1 Design and Implementation Issues

6.1.1 A guard cannot test for the absence of a failure it never produces

Three defects shared one shape, described in Section 4.3.3, and it is the single most transferable finding of the project. Each was a check written against silence when the real failure was confident wrongness. The lesson is to test for the shape of a right answer rather than for the absence of a wrong one.

6.1.2 Constants standing in for measurements

Two panes sized themselves as `calc(100vh - 11rem)`, with a comment estimating the height of the page heading. Removing the heading made the constant wrong, silently, and the correction had to be re-derived per screen. Replacing it with a layout that takes the height it is given failed on the first attempt, because the property used resolves against the nearest enclosing flexible container and an unclassed wrapper sat between the two. The general form: a guessed constant fails quietly, and replacing a mechanism requires knowing what it sits inside.

6.1.3 Measuring on material one has chosen

The extractor scored 100% on the authors' own test set and 62.5% on a random sample. Nothing about the development set was dishonest; it was drawn from a single field, and a template is a convention of a field.

6.1.4 Replacing a running interface costs double

The interface was rebuilt as a parallel tree while the original kept shipping. Every fix made during that period was made twice, and copying between the trees was not available, because the replacement names its colour tokens the way the original named its sizes. The mitigation, running both through the same CI matrix, and parameterising the contract test over both, was necessary precisely because the tree nobody is looking at is the one that rots.

6.2 Advantages and Limitations

6.2.1 Advantages

- No document, prompt or embedding leaves the machine, which makes the system usable on material that cannot be uploaded.

- No per-query cost and no account.
- Fully functional offline, including typesetting.
- Metadata accurate enough to label a paper and populate a citation: 92.9% of titles on unseen papers.
- Sized to the machine, so a consumer laptop is the target rather than the degraded case.
- Distributed and verified as installers on three platforms, with the packaged application itself under test.

6.2.2 Limitations

- **Model capability.** A 0.5B model produces malformed structured output on the harder tasks; those route to a 1.5B when one is present and degrade when it is not.
- **Scanned documents.** No text layer means nothing to read. OCR is not shipped.
- **Extraction tail.** Mathematics inside a title, publisher cover pages, and submissions with no identifier still defeat the rules.
- **Retrieval quality.** Measured on one query, the top-ranked passage scored 0.212 where a good match on this embedding model is 0.45–0.70, and a cover page outranked the passage containing the answer. The probable cause is chunk size diluting the signal. This is unresolved and is stated rather than omitted.
- **Citations.** Entries can be generated from extracted metadata but will be incomplete: venue, volume and pages are not recoverable from page one.
- **No collaboration or synchronisation,** which follows from the premise.

6.3 Future Scope

1. **A layout classifier for metadata.** The remaining extraction failures are a long tail of typesetting conventions, and each currently costs another rule. The established answer is a small classifier over the same features the rules already read: font size, position, capitalisation, row order, which is what GROBID does with a conditional random field at roughly 90% F1. The research interest is not the classifier but the constraint: GROBID needs a JVM and a model server, and the neural alternative needs a GPU. *Nobody optimises this task for a CPU-only machine with no network*, which is the constraint this application is built under, and measuring how close a fixed budget gets to those ceilings is publishable either way.
2. **Remembering corrections.** When a user fixes a wrong title, store the page’s layout fingerprint against the correction. Papers arrive in clusters from the same venues, so one correction pays for itself repeatedly. This is adaptation without gradients: inspectable, reversible, and free of the forgetting that on-device fine-tuning of a shared model would risk.
3. **Source-to-output synchronisation.** The compiler already emits a SyncTeX map on every build and nothing reads it. Click-through in both directions is available for the cost of parsing a file already on disk.

4. **Fine-tuned task models.** A model trained on instruction-to-markup pairs, for which data is already collected passively during use.
5. **Retrieval quality.** The 0.212 result deserves its own investigation with a proper query set.

REFERENCES

1. Vaswani, A., Shazeer, N., Parmar, N., *et al.* “Attention Is All You Need.” *Advances in Neural Information Processing Systems*, 2017. arXiv:1706.03762.
2. Reimers, N. and Gurevych, I. “Sentence-BERT: Sentence Embeddings using Siamese BERT-Networks.” *EMNLP*, 2019. arXiv:1908.10084.
3. Lewis, P., Perez, E., Piktus, A., *et al.* “Retrieval-Augmented Generation for Knowledge-Intensive NLP Tasks.” *NeurIPS*, 2020. arXiv:2005.11401.
4. Lopez, P. “GROBID: Combining Automatic Bibliographic Data Recognition and Terminology Recognition.” *ECDL*, 2009.
5. Blecher, L., Cucurull, G., Scialom, T. and Stojnic, R. “Nougat: Neural Optical Understanding for Academic Documents.” 2023. arXiv:2308.13418.
6. Yang, A., Yang, B., Hui, B., *et al.* “Qwen2.5 Technical Report.” 2024. arXiv:2412.15115.
7. Gerganov, G. *llama.cpp: LLM inference in C/C++*. 2023. <https://github.com/ggerganov/llama.cpp>
8. Biggs, N. *Tectonic: A modernized, complete, self-contained T_EX engine*. <https://tectonic-typesetting.github.io>
9. Tauri Contributors. *Tauri: Build smaller, faster, more secure desktop applications*. <https://tauri.app>
10. Ramírez, S. *FastAPI*. <https://fastapi.tiangolo.com>
11. Biham, E. and Dunkelman, O. “Argon2: the memory-hard function for password hashing.” *IETF RFC 9106*, 2021.
12. Adams, R. P. and MacKay, D. J. C. “Bayesian Online Changepoint Detection.” 2007. arXiv:0710.3742.